

Hands-on: Clinical NLP with Python

Clinical NLP Pipeline Example

```
# Install: pip install scispacy import spacy import scispacy from scispacy.linker import EntityLinker # Load clinical model nlp = spacy.load("en_core_sci_md") nlp.add_pipe("scispacy_linker", config={"resolve_abbreviations": True}) # Process clinical text text = "Patient prescribed metformin 500mg for T2DM. HbA1c was 7.2%" doc = nlp(text) # Extract entities for ent in doc.ents: print(f" {ent.text} > {ent.label_}") if ent._.kb_ents: cui = ent._.kb_ents[0][0] # UMLS CUI print(f" UMLS: {cui}")
```

```
# BioBERT for NER using Transformers from transformers import AutoTokenizer, AutoModelForTokenClassification import torch # Load BioBERT model tokenizer = AutoTokenizer.from_pretrained("dmis-lab/biobert-v1.1") model = AutoModelForTokenClassification.from_pretrained( "dmis-lab/biobert-v1.1" ) # Tokenize and predict inputs = tokenizer(text, return_tensors="pt") outputs = model(**inputs) predictions = torch.argmax(outputs.logits, dim=2)
```

NLP Pipeline Flow



- pip install scispacy
- Biomedical NER models
- UMLS entity linking



- Medical Concept Annotation
- Unsupervised learning
- Active learning interface

- Abbreviation detection
- Negation with NegEx

- Context detection
- SNOMED/UMLS linking



BioBERT/ClinicalBERT

- Pretrained on PubMed/MIMIC
- Fine-tuning for NER
- Relation extraction
- Question answering
- Hugging Face integration



Evaluation Metrics

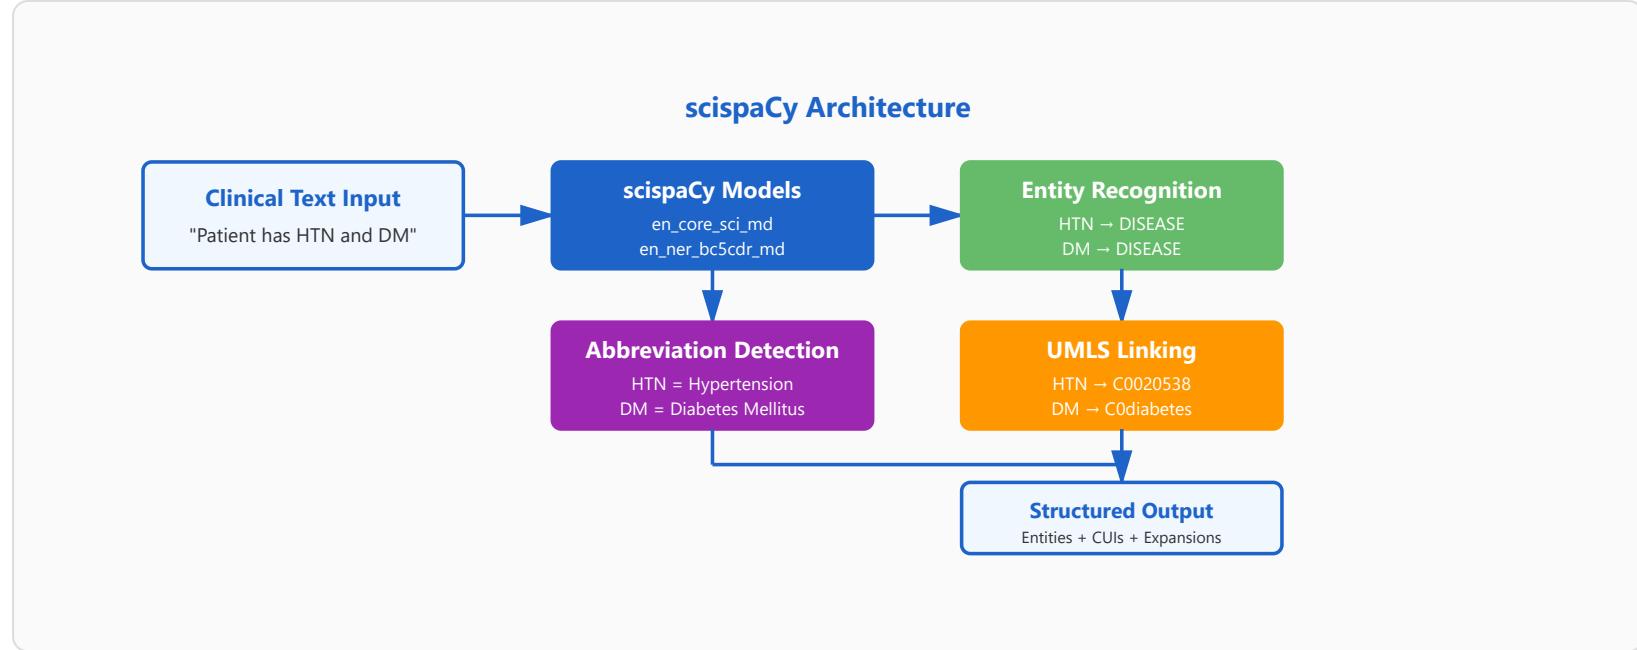
- Precision: correct/predicted
- Recall: correct/actual
- F1 score: harmonic mean
- Entity vs token level
- Cross-validation

Detailed Framework Explanations



scispacy: Scientific spaCy for Biomedical Text

scispacy is a Python package built on top of spaCy, specifically designed for processing biomedical and clinical text. It provides specialized models trained on scientific literature and includes tools for named entity recognition, entity linking to standardized vocabularies like UMLS, and abbreviation detection.



🎯 Pretrained Models

Multiple specialized models trained on biomedical literature: en_core_sci_sm, en_core_sci_md, en_core_sci_lg for general scientific text, and en_ner_bc5cdr_md for chemical and disease entities.

🔗 Entity Linking

Automatic linking to UMLS, MeSH, RxNorm, Gene Ontology, and other knowledge bases. Maps recognized entities to standardized concept unique identifiers (CUIs).

📝 Abbreviation Detection

Detects and expands medical abbreviations using a specialized component. Handles both explicit definitions and implicit abbreviations in clinical text.

🚫 Negation Detection

Integration with NegEx algorithm to detect negated medical concepts. Essential for understanding "patient denies chest pain" vs "patient has chest pain".

💡 Use Case Example

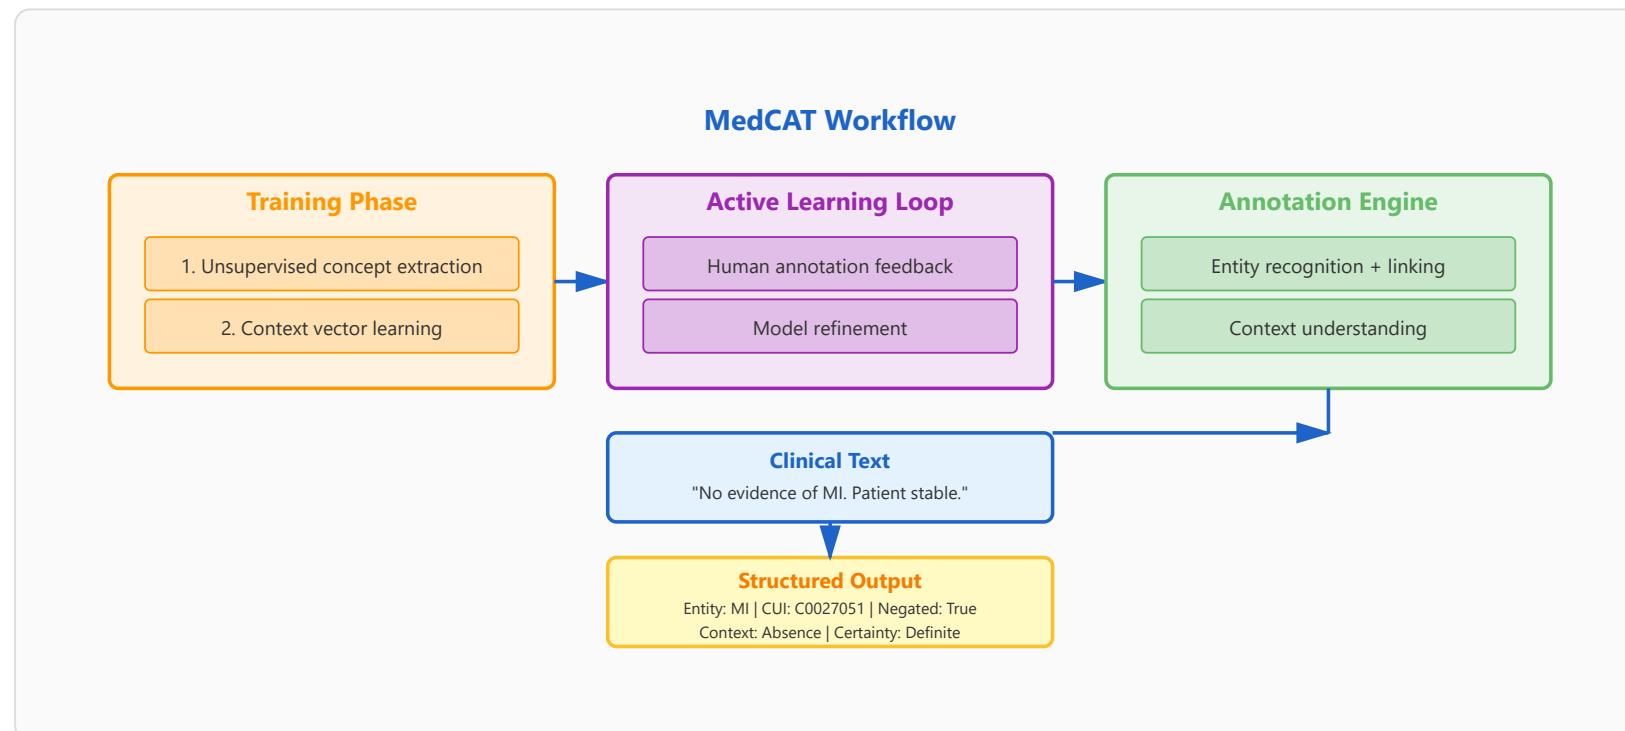
Clinical Note Processing: A hospital wants to extract structured information from discharge summaries. scispacy can identify diseases (ICD codes), medications (RxNorm), procedures, and link them to standardized vocabularies for downstream analytics and quality reporting.

```
# Complete scispacy workflow
import spacy
from scispacy.abbreviation import AbbreviationDetector
from scispacy.linker import EntityLinker
nlp = spacy.load("en_core_sci_sm")
nlp.add_pipe("abbreviation_detector")
nlp.add_pipe("scispacy_linker", config={"resolve_abbreviations": True, "linker_name": "umls"})
text = "Patient with CHF presents with dyspnea. Started on furosemide."
doc = nlp(text)
for ent in doc.ents:
    print(f"Entity: {ent.text}, Label: {ent.label_}")
    for umls_ent in ent._.kb_ents:
        print(f"UMLS CUI: {umls_ent[0]}, Score: {umls_ent[1]}")
```



MedCAT: Medical Concept Annotation Tool

MedCAT is an advanced NLP framework designed for medical concept annotation in clinical text. It uses unsupervised learning to create medical concept models and provides an active learning interface for continuous model improvement. MedCAT excels at handling context, negation, and temporal information.



Unsupervised Learning



Context Detection

MedCAT can learn medical concepts from unlabeled text using context-based embeddings. It builds vocabulary and concept relationships automatically from large clinical corpora without manual annotation.

Advanced context analysis detects negation, temporality (past/present/future), experiencer (patient/family), and certainty. Understands "no history of diabetes" vs "diagnosed with diabetes".

Active Learning Interface

MedCAT Trainer provides a web-based annotation interface where clinicians can correct and improve models. The system learns from corrections and improves over time.

Multi-Ontology Support

Links entities to SNOMED CT, UMLS, ICD-10, and custom ontologies. Supports hierarchical relationships and allows navigation of medical concept taxonomies.

Use Case Example

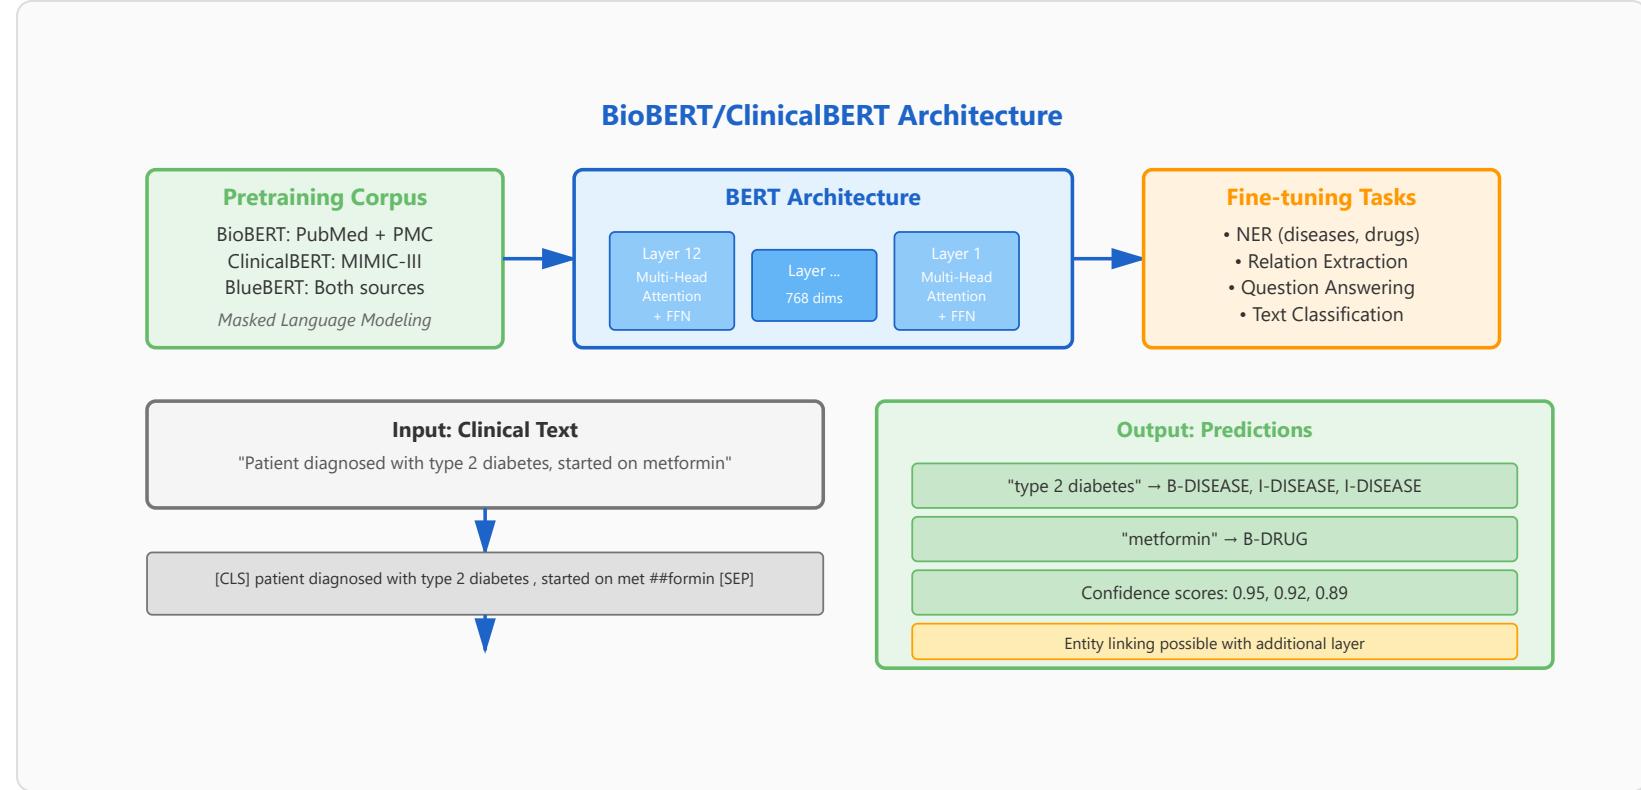
Clinical Decision Support: A healthcare system deploys MedCAT to extract structured data from clinical notes in real-time. The system identifies diagnoses, their context (confirmed, suspected, ruled out), and links them to SNOMED CT codes for clinical decision support and automated quality measures.

```
# MedCAT implementation example
from medcat.cat import CAT # Load pretrained model
cat = CAT.load_model_pack("path/to/model_pack.zip") # Process clinical text
text = """Patient denies chest pain. History of hypertension, well controlled on lisinopril. No signs of heart failure."""
entities = cat.get_entities(text)
for entity in entities['entities'].values():
    print(f"\nConcept: {entity['pretty_name']}") 
    print(f"SNOMED: {entity['cui']}") 
    print(f"Negated: {entity['meta_anno']['Negation']}") 
    print(f"Temporality: {entity['meta_anno']['Temporality']}")
```



BioBERT & ClinicalBERT: Transformer Models for Healthcare

BioBERT (Biomedical BERT) and **ClinicalBERT** are transformer-based language models pretrained on biomedical literature and clinical notes respectively. They leverage the BERT architecture with domain-specific pretraining to achieve state-of-the-art performance on various biomedical NLP tasks including named entity recognition, relation extraction, and question answering.



📚 Domain Pretraining

BioBERT trained on 4.5B words from PubMed abstracts and PMC full-text articles. ClinicalBERT trained on MIMIC-III clinical notes. This domain-specific pretraining provides deep understanding of medical language and terminology.

🎯 State-of-the-Art Performance

Achieves best results on BC5CDR (chemical/disease NER), ChemProt (relation extraction), and BioASQ (question answering). F1 scores typically 85-90% on biomedical NER tasks.

🔧 Easy Fine-tuning

Integrated with Hugging Face Transformers library. Simple fine-tuning API with just a few lines of code. Supports transfer learning for new tasks with limited labeled data (few-shot learning).

🌐 Multiple Variants

Various sizes available: BioBERT-Base (110M params), BioBERT-Large (340M params), BlueBERT (PubMed + MIMIC), PubMedBERT (from scratch). Choose based on accuracy vs speed requirements.

💡 Use Case Example

Clinical Trial Matching: A research institution uses BioBERT to extract patient conditions and medications from EHRs, then matches patients to relevant clinical trials based on inclusion/exclusion criteria. The model identifies complex medical concepts and their relationships with high accuracy.

```
# BioBERT for Named Entity Recognition from transformers import ( AutoTokenizer, AutoModelForTokenClassification, pipeline ) #
Load BioBERT NER model model_name = "dmis-lab/biobert-base-cased-v1.2" tokenizer = AutoTokenizer.from_pretrained(model_name) model
= AutoModelForTokenClassification.from_pretrained( "alvaroalon2/biobert_diseases_ner" ) # Create NER pipeline ner_pipeline =
pipeline( "ner", model=model, tokenizer=tokenizer, aggregation_strategy="simple" ) # Process text text = """Patient presents with
acute myocardial infarction. History of type 2 diabetes mellitus and hypertension. Started on aspirin and metformin.""" entities =
ner_pipeline(text) for entity in entities: print(f"{entity['word']} | {entity['entity_group']} | Score: {entity['score']:.3f}")
```



Evaluation Metrics for Clinical NLP

Evaluation metrics are essential for measuring the performance of NLP models. In clinical NLP, we need to carefully assess how well models identify medical entities, link them to correct concepts, and maintain high accuracy across diverse clinical texts. Understanding precision, recall, and F1 score helps us compare models and ensure they meet clinical requirements.

Evaluation Metrics Framework

Confusion Matrix

		Predicted	
		Positive	Negative
Actual	Positive	TP True Positive Correctly identified	FP False Pos Incorrect
	Negative	FN False Neg Missed	TN True Neg Correct

Metric Formulas

Precision:
 $TP / (TP + FP)$

"How many predicted entities are correct?"

Recall:
 $TP / (TP + FN)$

"How many actual entities were found?"

F1 Score:
 $2 \times (Precision \times Recall) / (Precision + Recall)$

"Harmonic mean of both metrics"

Accuracy:

$(TP + TN) / (TP + TN + FP + FN)$ "Overall correctness"

Specificity:

$TN / (TN + FP)$ "True negative rate"

Example: Disease Entity Recognition

Ground Truth: ["diabetes", "hypertension", "heart failure"]

Predictions: ["diabetes", "hypertension", "cardiac arrest"]

TP = 2 (diabetes, hypertension)
FP = 1 (cardiac arrest)

FN = 1 (heart failure missed)
TN = N/A (open-ended task)

Precision = $2/3 = 0.67$ (67%)
Recall = $2/3 = 0.67$ (67%)
F1 Score = 0.67

Precision

Measures correctness of predictions. High precision means few false positives. Critical in clinical settings to avoid incorrect diagnoses or treatment recommendations.

Example: Of 100 predicted diseases, how many are actually correct?

Recall (Sensitivity)

Measures completeness of predictions. High recall means few missed entities. Essential for patient safety to ensure no critical conditions are overlooked. Example: Of 100 actual diseases in text, how many did we find?

F1 Score

Harmonic mean balancing precision and recall. Single metric for model comparison. F1 of 0.85+ considered good for clinical NLP. Use F1-macro for class imbalance, F1-micro for overall performance.

Evaluation Levels

Token-level: Each token scored separately. **Entity-level:** Entire entity must match exactly (strict) or overlap partially (relaxed). **Span-level:** Boundaries must match. Clinical NLP typically uses strict entity-level evaluation.

Use Case Example

Model Selection for Deployment: A hospital compares three NER models: Model A (Precision: 0.95, Recall: 0.75, F1: 0.84), Model B (P: 0.85, R: 0.90, F1: 0.87), Model C (P: 0.88, R: 0.88, F1: 0.88). They choose Model B for a screening application where missing conditions is more dangerous than false positives, while Model A might be better for automated billing where precision matters more.

```
# Calculate evaluation metrics for NER from sklearn.metrics import precision_recall_fscore_support, classification_report import numpy as np # True labels and predictions (IOB format) y_true = ['O', 'B-DIS', 'I-DIS', 'O', 'B-DRUG', 'O'] y_pred = ['O', 'B-DIS', 'I-DIS', 'O', 'O', 'O'] # Calculate metrics precision, recall, f1, support = precision_recall_fscore_support( y_true, y_pred, average='weighted', zero_division=0 ) print(f"Precision: {precision:.3f}") print(f"Recall: {recall:.3f}") print(f"F1 Score: {f1:.3f}") # Detailed report by entity type print("\nClassification Report:") print(classification_report(y_true, y_pred)) # Entity-level evaluation (exact match) def entity_level_f1(true_entities, pred_entities): tp = len(true_entities.intersection(pred_entities)) fp = len(pred_entities - true_entities) fn = len(true_entities - pred_entities) precision = tp / (tp + fp) if (tp + fp) > 0 else 0 recall = tp / (tp + fn) if (tp + fn) > 0 else 0 f1 = 2 * precision * recall / (precision + recall) if (precision + recall) > 0 else 0 return precision, recall, f1 # Example usage true_ents = {'diabetes': 10, 18}, ('hypertension', 25, 37)} pred_ents = {'diabetes', 10, 18}, ('heart failure', 40, 53)} p, r, f = entity_level_f1(true_ents, pred_ents) print(f"\nEntity-level - Precision: {p:.2f}, Recall: {r:.2f}, F1: {f:.2f}")
```

Clinical NLP Best Practices

Always validate models on diverse clinical datasets • Consider domain adaptation for your specific use case • Monitor performance in production
• Ensure compliance with healthcare regulations (HIPAA, GDPR) • Involve clinical experts in evaluation • Test for bias across patient demographics